



COURSE SLIDES

AI for HR

Course Code: C820

Conducted by Tertiary Infotech Academy Pte Ltd · UEN 201200696W

Version v1.0 · 29 July 2026

© 2026 Tertiary Infotech Academy Pte Ltd. All rights reserved. · www.tertiarycourses.com.sg



COURSE ADMINISTRATION

Welcome & Housekeeping

Course Trainer Profiles

1

Assigned Tertiary Infotech Academy Trainer

HR-technology and agentic-AI facilitator who guides human-owned workflow design, evidence-led prompting, fair recruitment support, policy-grounded employee service and practical deployment controls.

Let's Know Each Other

1 Your name and organisation / role.

2 Your experience with this subject (if any).

3 What you want to be able to do after this course.

Ground Rules

1

Set your mobile phone to silent mode.

2

Participate actively — no question is too small.

3

Mutual respect: agree to disagree.

4

One conversation at a time.

5

Be punctual; return from breaks on time.

6

Learn by doing — try every hands-on lab.

Access Course Materials

1

Open the C820 course page or the
trainer-supplied Drive folder

2

Open the courseware links for this
course

3

Download the Slides, Learner Guide
and Lesson Plan

4

Open the Activities folder for all 6
labs

5

Keep the files together as you work

Course page: <https://www.tertiarycourses.com.sg/ai-for-hr.html>

Lesson Plan — 1 Day, 7.5 instructional hours

Course sequence

- Day 1 — Prompt, Build and Govern - create a controlled HR-agent portfolio from approved evidence
 - Topic 1: Get Started with AI for HR (Labs 1–2)
 - Topic 2: Build HR AI Agents (Labs 3–4)
 - Topic 3: Deploy Agentic AI Across the HR Function (Labs 5–6)

Day format

- Daily timing
 - 9:30 am-6:30 pm (1-hour lunch; two 15-minute tea breaks; 7.5 instructional hours)
- Instructional time
 - 7.5 hours · 6 connected labs
- Learning arc
 - Concept → demonstrate → build → verify → discuss → recap

Learning Outcomes

1

Prompt & Ground

Use an approved AI assistant with a structured prompt, source boundary and human review method for everyday HR work.

2

Design HR Agents

Design bounded recruitment, onboarding, employee-support and communication agents with job-related criteria and explicit human decisions.

3

Automate & Verify

Build and test a multi-step HR workflow that connects approved data and tools while preserving exceptions, approvals and run evidence.

4

Govern & Roll Out

Evaluate privacy, fairness, security and operating risk, then produce a staged rollout plan with owners, metrics, monitoring and fallback.

How You'll Learn

1

Trainer demonstrates the concept

2

You build it yourself in the hands-on lab

3

Verify with the lab's 'Test it' step

4

Discuss and troubleshoot together

5

Recap against the learning outcomes



CORE CONCEPTS

The Human-Owned HR Agent

Four Boundaries Before Any Build

1

Purpose boundary

Define the HR outcome, affected people, non-goals and accountable owner.

2

Evidence boundary

Use necessary approved sources with version, owner and lineage.

3

Action boundary

Separate read, draft, recommend, approve and execute permissions.

4

Decision boundary

Keep consequential people decisions with an authorised human who can change the result.

One Controlled HR-Agent Run

1

Receive

Validate requester, purpose, case and permitted data.

2

Ground

Retrieve the current approved role, policy or workflow source.

3

Work

Apply fixed rules and bounded model tasks.

4

Review

Check evidence, fairness, privacy, tone and action authority.

5

Decide

A named human approves, edits, escalates or stops.

6

Record

Retain sources, checks, decision, action and final status.

COURSE OPERATING PRINCIPLE

Use AI to organise HR work - never to hide how a people decision was made.

Every material claim has a source, every exception has an owner and every consequential decision remains human-owned.

What You Will Build

Foundation

- HR use-case and control canvas
- grounded prompt and source pack

HR Agents

- recruitment evidence assistant
- onboarding and employee-support pack

Deployment

- multi-step workflow simulation
- governance test and rollout plan

How Every Lab Progresses

1

Open the approved synthetic Asteron People Operations checkpoint.

2

Apply a prompt contract, rubric or workflow rule to bounded inputs.

3

Separate source facts, model output, deterministic checks and human decisions.

4

Run the observable Test It checks and repair defects.

5

Save the evidence and rejoin checkpoint for the next lab.

Assistant, Automation or Agent?

1

Assistant

A person chooses each request and decides what to do with the answer.

2

Automation

Fixed rules execute a known sequence with predictable branches.

3

Agent

A model chooses among approved next actions inside explicit limits.

4

Design rule

Use the least autonomy that meets the HR need.

The HR Evidence Stack

1

Authoritative source

Current role, policy, record or approved process definition.

2

Deterministic layer

Required fields, dates, identifiers, formulas, thresholds and routing.

3

Model layer

Extraction, classification, drafting, critique and exception explanation.

4

Decision layer

Named human authority, reason, communication and controlled action.

The Five Human Review Questions

1

Evidence

Which approved source supports each material statement?

2

Fairness

Are criteria job-related, consistent and free of unsupported inference?

3

Privacy

Is each data field necessary and handled in an approved service?

4

Authority

May this system only draft, or may it take the proposed action?

5

Recourse

Can the affected person or operator question and correct the result?

TOPIC 01

Get Started with AI for HR

AI and agentic AI use cases - tools and platforms - effective prompting - approved HR data and documents

Key Concepts — Get Started with AI for HR

1

Assistant, automation, agent

Choose the least autonomous pattern that can complete the HR job reliably.

2

Human-owned purpose

Name the employee or candidate outcome, process owner and decision that remains human.

3

Prompt contract

State goal, context, criteria, sources, output, review and escalation before using a model.

4

Grounded evidence

Constrain material claims to approved policy, role and case sources and preserve their identifiers.

5

Data minimisation

Use synthetic or necessary authorised fields and keep confidential or sensitive data out of unapproved tools.

6

Observable quality

Define acceptance checks before comparing model outputs or platforms.

AI and Agentic AI in HR: Use Cases and Landscape

WHAT

- Generative AI produces or transforms language and other content from instructions and context. An AI assistant responds to a person, a deterministic automation follows fixed rules, and an agent lets a model choose among approved next steps and tools inside a bounded workflow.

WHY IT MATTERS

- HR work combines high-volume drafting with decisions that affect people. Separating assistance, automation and agency prevents a fluent tool from being given authority that belongs to an HR professional, hiring manager, data owner or employee.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — AI and Agentic AI in HR: Use Cases and Landscape

1

Define the HR outcome, affected people, process owner and authoritative sources.

2

Classify each step as fixed rule, model-supported judgement or human decision.

3

Grant only the data and tools required for the current step.

4

Set completion, uncertainty, exception, time and action limits.

5

Retain the source, output, checks, edits, decision and final status as run evidence.

In Practice — AI and Agentic AI in HR: Use Cases and Landscape

WORKED EXAMPLE

- Asteron Services wants faster onboarding support for new employees.
- A deterministic rule identifies the employee's country and start date; an assistant drafts a policy-grounded reply.
- A named HR owner reviews exceptions and any message about eligibility, conduct, pay or personal circumstances.

DECISION GUIDE

- **Use when**
 - The task contains repeated language work, unstructured documents or exceptions that fixed rules handle poorly.
 - A clear source boundary, observable checks and accountable reviewer can be defined.
- **Avoid when**
 - A template, search filter or spreadsheet rule solves the task more simply and predictably.
 - The system would make an adverse or irreversible people decision without meaningful human review.

Quality Lens — AI and Agentic AI in HR: Use Cases and Landscape

Bounded

- Purpose, sources, tools, limits, stop conditions and fallback are explicit.

Grounded

- Material statements trace to approved facts or rules.

Accountable

- A named role owns the final people decision and communication.

Popular AI Tools and Agent Platforms

WHAT

- HR teams can work with browser assistants such as ChatGPT or Claude, configurable workspace agents, low-code platforms such as Copilot Studio, or custom software. The options combine models, instructions, knowledge and tools but differ in data handling, integration, testing, identity, observability, cost and operating skill.

WHY IT MATTERS

- Choosing a product before defining the use case encourages feature-led adoption. A task-and-risk comparison makes the platform decision explainable and avoids connecting broad HR repositories to a prototype that needs only a synthetic policy extract.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Popular AI Tools and Agent Platforms

1

Start with the HR job, data class, actions, users, volume and acceptable failure mode.

2

Prototype a prompt with synthetic sources in an approved supervised assistant.

3

Move to a workspace agent when reusable instructions and governed knowledge are sufficient.

4

Use low-code flow tooling for triggers, connectors, branching, approvals and monitoring.

5

Use custom development only when identity, tools, evaluation, telemetry or isolation require it.

In Practice — Popular AI Tools and Agent Platforms

WORKED EXAMPLE

- A one-off job-description critique stays in a supervised assistant with a synthetic role brief.
- A reusable employee-policy helper uses approved knowledge and read-only access.
- An onboarding workflow uses a flow platform because it needs triggers, record lookups, tasks, approvals and run history.

DECISION GUIDE

- **Use when**
 - Comparing platforms against an already documented HR workflow and control matrix.
 - The organisation can confirm approved services, data locations, identities and support ownership.
- **Avoid when**
 - Selecting the tool with the most connectors or the highest apparent autonomy.
 - Allowing inherited user permissions to expose unrelated employee or candidate records.

Quality Lens — Popular AI Tools and Agent Platforms

Fit

- Workflow and risk needs drive the tool choice.

Control

- Identity, data, actions and environments can be restricted.

Visibility

- Versions, runs, failures, approvals and costs can be inspected.

Effective Prompting for HR Tasks

WHAT

- An HR prompt is a working contract. This course uses G-C-C-S-O-R: Goal, Context, Criteria, Sources, Output and Review. The contract separates stable instructions from case variables, defines the evidence boundary and makes the desired result testable.

WHY IT MATTERS

- Requests such as 'screen these resumes' or 'write an onboarding email' hide the selection basis, source truth and acceptance checks. A structured prompt reduces generic output and makes unsupported claims, missing facts and inappropriate inferences easier to detect.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Effective Prompting for HR Tasks

1

State one HR goal, intended reader and permitted decision support.

2

Provide only relevant context and label each approved source with a stable identifier.

3

Define job-related or policy-based criteria before asking for analysis.

4

Specify output fields for facts, source, uncertainty, recommendation and human action.

5

Require UNKNOWN, clarification or escalation when evidence is absent.

6

Review accuracy, fairness, privacy, rights, tone and action authority before use.

In Practice — Effective Prompting for HR Tasks

WORKED EXAMPLE

- Goal: draft a first-day email for a Singapore-based new employee.
- Sources: offer-confirmation facts and approved onboarding policy excerpts; personal medical or family information is excluded.
- Output: subject, message, source ledger and unresolved questions; any policy exception is marked ESCALATE TO HR.

DECISION GUIDE

- **Use when**
 - The same HR transformation or review task will recur with different approved inputs.
 - Success criteria can be checked against a source, rubric, schema or reviewer checklist.
- **Avoid when**
 - The purpose, authorised source or intended human decision is unclear.
 - The prompt would include real personal data in a service not approved for that data.

Quality Lens — Effective Prompting for HR Tasks

Specific

- Goal, scope, audience, criteria and format are explicit.

Evidence-led

- Sources are delimited and every material claim carries a reference.

Fail-safe

- Missing or conflicting evidence triggers a question or escalation.

Connecting AI to HR Data and Documents

WHAT

- A governed HR connection exposes a defined source through a limited interface. Its data contract records purpose, owner, fields, record grain, version, refresh time, permitted users, retention, quality checks and whether the connection may read or write.

WHY IT MATTERS

- Candidate and employee repositories contain personal and sometimes sensitive information. Broad access increases privacy, prompt-injection and accidental-action risk, while unclear versions allow an agent to quote a retired policy as if it were current.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Connecting AI to HR Data and Documents

1

Classify the use case and minimise rows, fields, periods and documents before connection.

2

Begin with synthetic or de-identified snapshots and read-only tools.

3

Record authoritative owner, effective date, version, jurisdiction and expiry for every policy.

4

Validate completeness, duplicates, field meanings, access and document status before retrieval.

5

Return source identifiers and timestamps with retrieved content.

6

Separate retrieval from any write, send or status-change tool and place approvals before actions.

In Practice — Connecting AI to HR Data and Documents

WORKED EXAMPLE

- The employee helper can retrieve three approved policy excerpts by Policy_ID and Effective_Date.
- It cannot browse payroll, performance or medical folders and has no send permission.
- A conflicting or expired excerpt stops the run and creates a question for the policy owner.

DECISION GUIDE

- **Use when**
 - A repeatable workflow needs approved structured records or policy documents.
 - A data owner can define purpose, quality, access, retention and lineage.
- **Avoid when**
 - The connector exposes an entire HR drive or inherited administrator privileges.
 - The source has no owner, effective date, stable identifier or current-version rule.

Quality Lens — Connecting AI to HR Data and Documents

Minimal

- Only the necessary records and fields are available.

Current

- Version, effective date and authoritative owner are visible.

Traceable

- Each output points to the exact source retrieved for that run.

Hands-On Labs — Get Started with AI for HR

Lab 1

- Frame the Asteron HR Agent and Write the Prompt Contract

Lab 2

- Build the Grounded HR Source and Employee FAQ Pack

Frame the Asteron HR Agent and Write the Prompt Contract

LAB 1

You will inspect the synthetic Asteron People Operations brief, compare an assistant, fixed automation and agent pattern, and choose a bounded onboarding-support use case. You will map its purpose, evidence, action and decision boundaries before running a G-C-C-S-O-R prompt and reviewing the result.

You'll build

01-foundation/hr-agent-foundation.md containing the selected pattern, four-boundary canvas, first prompt, initial and refined outputs, claim ledger, human decision and stop conditions.

1 Choose the use case

2 Map four boundaries

3 Write G-C-C-S-O-R

4 Run and review

Tools: Text editor - approved AI assistant - asteron-people-operations-brief.md

Frame the Asteron HR Agent and Write the Prompt Contract

Test it

Open 01-foundation/hr-agent-foundation.md. It must name one pattern and justify it; contain all four boundaries; contain one explicit completion, uncertainty, time, duplicate and unsafe-action condition; contain one complete G-C-C-S-O-R prompt, initial and refined outputs, a version-tagged claim ledger with no unsupported statement marked SUPPORTED, and one human decision.

Build the Grounded HR Source and Employee FAQ Pack

LAB 2

You will inspect synthetic onboarding policy excerpts, identify the current source for each policy topic and define a read-only retrieval contract. You will then produce six answer-or-handoff records, cite the exact Policy_ID and effective date for answers, and route cases that cannot be answered safely.

You'll build

01-foundation/hr-source-register.csv and 01-foundation/employee-faq-pack.md with current-version rules, six answer-or-handoff records, source citations, unknowns, escalation routes and a retrieval test log.

1 Validate sources

2 Set retrieval rules

3 Draft answers

4 Verify lineage

Tools: Spreadsheet - text editor - approved AI assistant - approved-policy-excerpts.md - employee-questions.csv

Build the Grounded HR Source and Employee FAQ Pack

Test it

The source register must contain every supplied policy version, exactly one CURRENT row per topic, and no blank owner or effective date. The FAQ pack must cover Q-001 to Q-006, cite Policy_ID and Effective_Date for every ANSWER, contain no conclusion for CLARIFY or ESCALATE cases, and show all four retrieval tests PASS.

Recap — Get Started with AI for HR

1

You can now: LO1: use an approved AI assistant with a structured prompt, source boundary and human review method

2

You can now: LO1: connect an HR task to current approved sources while preserving source identity, uncertainty and human review

TOPIC 02

Build HR AI Agents

agent workflow design - recruitment decision support - onboarding and employee support - HR content and communications

Key Concepts — Build HR AI Agents

1

Workflow contract

Map trigger, sources, tools, branches, decisions, evidence and fallback before configuration.

2

Job-related rubric

Translate role requirements into consistent observable criteria before reviewing candidates.

3

Decision support

The agent organises evidence and questions; an authorised person owns the hiring decision.

4

Policy-grounded support

Answer from current approved policy and escalate personal, ambiguous or high-impact cases.

5

Human communication

Use AI for options and editing while a responsible sender verifies facts, tone and audience.

6

Verification

Test normal, boundary, missing-data, conflicting-source and unsafe-action cases.

Designing Agentic AI Workflows for HR

WHAT

- An HR agent workflow is a bounded sequence in which a model may choose among approved read, transform, classify, ask and draft steps. The design contract defines triggers, inputs, state, tools, branches, human decisions, evidence and stop conditions before a platform is configured.

WHY IT MATTERS

- A demonstration often shows only the happy path. HR operations also contain missing documents, policy conflicts, candidate questions, protected information and exceptions that need a clear owner rather than an improvised model response.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Designing Agentic AI Workflows for HR

1

Write the outcome and non-goals, then draw the current human workflow.

2

Classify each step as deterministic rule, model task, human decision or external action.

3

Define tool inputs, outputs, permissions, timeouts and idempotency keys.

4

Add escalation for uncertainty, conflicting sources and high-impact cases.

5

Define completion evidence and a manual fallback before the first prototype.

In Practice — Designing Agentic AI Workflows for HR

WORKED EXAMPLE

- Trigger: an approved new-hire record is created with a start date and location.
- The flow checks required fields, retrieves current policy, drafts tasks and asks HR to approve the plan.
- Only after approval are tasks created; duplicate Employee_ID and Start_Date runs are blocked.

DECISION GUIDE

- **Use when**
 - The process is multi-step, has repeated exceptions and can be represented with explicit states.
 - Owners can specify allowed actions and what evidence proves completion.
- **Avoid when**
 - The desired outcome changes from case to case without a stable owner or policy basis.
 - The agent would need unrestricted access or an irreversible action without an approval boundary.

Quality Lens — Designing Agentic AI Workflows for HR

Complete

- Trigger, inputs, states, branches, owners and finish conditions are mapped.

Least privilege

- Every tool exposes only the fields and actions required.

Recoverable

- Retries, duplicates, failures and manual fallback are designed.

Recruitment and Candidate Screening Agents

WHAT

- A recruitment agent can extract role-relevant evidence, compare it with a pre-approved rubric, identify missing information and prepare consistent interview questions. It is decision support: it must not infer protected traits or replace the accountable hiring decision.

WHY IT MATTERS

- Unstructured resume review is vulnerable to inconsistent criteria and unsupported assumptions. A transparent rubric tied to actual job requirements improves consistency, while source citations, deferral and human review expose uncertainty instead of hiding it.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Recruitment and Candidate Screening Agents

1

Complete job analysis and approve skills, experience and evidence anchors before candidate review.

2

Remove non-job-related fields and instruct the system not to infer protected characteristics.

3

Extract evidence with resume line or record references; mark missing facts UNKNOWN.

4

Apply the same anchored rubric to every candidate for the same role.

5

Compare proposed and final human scores and record adjustments with reasons.

6

Monitor outcomes, error patterns and subgroup effects where lawful and appropriate.

In Practice — Recruitment and Candidate Screening Agents

WORKED EXAMPLE

- A role requires service-case handling, spreadsheet reporting and stakeholder communication.
- The agent cites evidence for each criterion and writes UNKNOWN when a resume lacks sufficient detail.
- The hiring panel reviews all evidence, corrects one over-score and owns the shortlist and interview plan.

DECISION GUIDE

- **Use when**
 - Criteria are job-related, documented and applied consistently across the candidate set.
 - The agent can cite evidence and defer uncertain cases to trained reviewers.
- **Avoid when**
 - Inferring age, disability, family status, ethnicity, health, personality or other sensitive traits.
 - Automatically rejecting a candidate or changing criteria after seeing candidate identities.

Quality Lens — Recruitment and Candidate Screening Agents

Relevant

- Every criterion is justified by the job analysis.

Consistent

- The same rubric and evidence rules apply to all candidates.

Reviewable

- Sources, unknowns, scores, changes and final human decision are retained.

Onboarding and Employee Support Agents

WHAT

- An onboarding or employee-support agent retrieves approved policy, gathers only necessary case facts, prepares tasks or a draft answer and routes exceptions to the correct owner. It should show the source and effective date and make escalation easy.

WHY IT MATTERS

- Fast answers improve the employee experience only when they are current and appropriate for the employee's context. A confident response from an expired policy can create more work and harm than a clear handoff to HR.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Onboarding and Employee Support Agents

1

Create an intent list and define which questions the agent may answer, draft or only route.

2

Retrieve current policy by stable identifier, jurisdiction and effective date.

3

Ask only for the minimum case facts needed to choose an approved route.

4

Return answer, source, confidence limits, next action and escalation path.

5

Separate support content from personal case decisions and employee-relations matters.

6

Log unanswered questions to improve the knowledge base through a controlled process.

In Practice — Onboarding and Employee Support Agents

WORKED EXAMPLE

- A new employee asks when to complete the equipment form and where to find it.
- The agent cites ONB-02, provides the current deadline and link, and labels the response as general guidance.
- A question about a personal leave exception is routed to an HR partner without the agent deciding eligibility.

DECISION GUIDE

- **Use when**
 - Approved policy and process sources are current, versioned and owned.
 - The support boundary and escalation routes are clear to employees and operators.
- **Avoid when**
 - A personal case requires judgement, accommodation, investigation, discipline or legal interpretation.
 - The knowledge source contains conflicting or expired versions without an authoritative rule.

Quality Lens — Onboarding and Employee Support Agents

Current

- Every answer shows the source and effective date.

Proportionate

- Only necessary case information is requested.

Escalatable

- High-impact and unresolved questions reach a named human route.

Drafting HR Content and Communications with AI

WHAT

- AI can generate and transform job descriptions, interview guides, onboarding messages, policy summaries and employee updates. The responsible sender supplies approved facts, reader context, tone and channel constraints, then verifies and edits before release.

WHY IT MATTERS

- HR language can shape access, trust and employee action. Generic or exclusionary wording, invented benefits and an incorrect deadline can cause real consequences even when the message sounds polished.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Drafting HR Content and Communications with AI

1

Name the audience, communication purpose, required action and approved source facts.

2

Specify plain-language, accessibility, inclusive-language and channel requirements.

3

Request options or critique rather than accepting the first draft.

4

Use a claim ledger to check dates, policy statements, links and promises.

5

Run a fairness, privacy, tone and action review and retain the final human edit.

In Practice — Drafting HR Content and Communications with AI

WORKED EXAMPLE

- An agent drafts a role advertisement from an approved job analysis and removes non-job-related preferences.
- It produces a first-day email from current onboarding policy and flags a missing equipment-contact name.
- The HR owner supplies the missing fact, edits the tone and approves the final versions.

DECISION GUIDE

- **Use when**
 - The communication has an accountable sender and a complete approved fact base.
 - The output can be checked against job, policy, brand and accessibility criteria.
- **Avoid when**
 - Sending high-impact or personalised messages automatically from an unreviewed draft.
 - Creating promises, role requirements or policy interpretations beyond the supplied source.

Quality Lens — Drafting HR Content and Communications with AI

Accurate

- Every date, promise, requirement and link is verified.

Inclusive

- Language is job-related, respectful and accessible.

Owned

- The named sender reviews the final message and intended audience.

Hands-On Labs — Build HR AI Agents

Lab 3

- Build the Fair Recruitment Evidence Assistant

Lab 4

- Build the Onboarding and Employee Support Agent Pack

Build the Fair Recruitment Evidence Assistant

LAB 3

You will turn an approved Asteron role profile into a transparent screening rubric, remove non-job-related fields from four synthetic candidate records and ask an AI assistant to cite evidence for each criterion. A human reviewer will correct proposed scores, record reasons and prepare interview questions without automatically rejecting or selecting anyone.

You'll build

02-hr-agents/recruitment-evidence-register.csv and 02-hr-agents/recruitment-review.md with criterion anchors, source evidence, unknowns, proposed and final scores, reviewer changes and human next-step decisions.

1 Approve criteria

2 Minimise inputs

3 Extract evidence

4 Human review

Tools: Spreadsheet - text editor - approved AI assistant - role-profile.md - synthetic-candidates.csv

Build the Fair Recruitment Evidence Assistant

Test it

The register must contain four Candidate_ID rows, three proposed and final components per row, exact evidence references, visible unknowns, totals from 0 to 6, and an Adjustment_Reason for every changed score. The review file must contain criterion anchors, six structured questions and one human-owned next step per candidate; it must contain no automated selection, ranking, rejection or protected-trait inference.

Build the Onboarding and Employee Support Agent Pack

LAB 4

You will combine the current source register from Lab 2 with six synthetic onboarding and employee-support cases. The agent will classify each case as ANSWER, CLARIFY or ESCALATE, draft a first-week plan and employee messages, and retain sources, limitations and a human release decision.

You'll build

02-hr-agents/onboarding-support-pack.md and **02-hr-agents/support-routing-register.csv** containing a first-week plan, six case routes, grounded drafts, source ledger, privacy check and human release decisions.

1 Classify cases

2 Retrieve policy

3 Draft two packs

4 Release review

Tools: Spreadsheet - text editor - approved AI assistant - onboarding-support-cases.csv - Lab 2 source register

Build the Onboarding and Employee Support Agent Pack

Test it

The routing register must contain six cases, necessary and omitted fields, one route, human owner and release decision per row. The support pack must include an ONB-ACT-01-grounded five-day plan for OS-001, six grounded drafts or handoffs, seven gate results and Final_Aproved_Text for every released case, with Policy_ID and Effective_Date for each ANSWER. All three boundary tests must PASS, and no message may be sent or personal exception decided.

Recap — Build HR AI Agents

1

You can now: LO2: design a recruitment decision-support agent with job-related criteria, source evidence and explicit human decisions

2

You can now: LO2: design onboarding and employee-support agents that use current policy, minimum necessary data and human communication review

TOPIC 03

Deploy Agentic AI Across the HR Function

multi-step automation - HR system integration - privacy, fairness and human oversight - staged organisational rollout

Key Concepts — Deploy Agentic AI Across the HR Function

1

Control-first automation

Design deterministic checks, model tasks and human gates as separate visible steps.

2

Integration contract

Specify identity, source, schema, permission, failure, retry and evidence for each connector.

3

Action boundary

Draft and recommend before any write, send, status change or people decision.

4

Risk-tiered oversight

Increase human involvement with uncertainty, sensitivity, impact and reversibility.

5

Evaluation set

Test normal, boundary, missing, conflicting, unsafe and fairness-relevant cases before pilot.

6

Staged rollout

Move from synthetic sandbox to shadow, read-only pilot and controlled operation with fallback.

Automating Multi-Step HR Workflows with AI Agents

WHAT

- A multi-step HR workflow coordinates triggers, deterministic rules, model tasks, human decisions and controlled actions over a visible state. The system completes only when required checks, approvals and evidence are present.

WHY IT MATTERS

- Automation can reduce handoffs but also repeats mistakes at speed. Separating deterministic truth, model-supported judgement and human authority makes the process easier to test, stop and recover.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Automating Multi-Step HR Workflows with AI Agents

1

Define states such as RECEIVED, VALIDATED, DRAFTED, NEEDS_REVIEW, APPROVED, ACTIONED and CLOSED.

2

Use deterministic checks for required fields, dates, identifiers, duplicates and routing.

3

Use the model only for bounded extraction, classification, drafting or critique.

4

Place approval before external messages, record changes and high-impact decisions.

5

Make actions idempotent and log state transitions, evidence and errors.

6

Provide timeout, retry, exception queue and manual completion routes.

In Practice — Automating Multi-Step HR Workflows with AI Agents

WORKED EXAMPLE

- An onboarding case enters RECEIVED and fails if Employee_ID, Start_Date or Location is absent.
- The model drafts the plan from approved policy; HR reviews exceptions and approves task creation.
- A repeated event with the same case key returns ALREADY_PROCESSED instead of creating duplicate tasks.

DECISION GUIDE

- **Use when**
 - The process has stable states, frequent handoffs and observable completion conditions.
 - Owners can define action permissions, approvals, exception service levels and fallback.
- **Avoid when**
 - A workflow is still changing daily and no stable current process exists.
 - Failures cannot be detected, reversed or completed through a manual route.

Quality Lens — Automating Multi-Step HR Workflows with AI Agents

Deterministic

- Rules own identifiers, dates, completeness, duplicates and thresholds.

Bounded

- Model tasks have explicit inputs, outputs and escalation.

Recoverable

- Every failure has an owner, retry rule and manual path.

Integrating Agents with HR Systems and Tools

WHAT

- An integration contract describes one tool call: authenticated identity, permitted operation, input schema, output schema, data scope, validation, timeout, retry, error handling, evidence and owner. Read, calculate, draft and action tools should remain distinct.

WHY IT MATTERS

- A connector inherits technical reach that may exceed the business use case. Explicit contracts reduce over-broad access and make it possible to test whether a run read the right record, used the right version and changed only what an authorised person approved.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Integrating Agents with HR Systems and Tools

1

Use a service or user identity with least privilege and a named owner.

2

Whitelist operations, fields, destinations and case types; deny everything else.

3

Validate input and output schemas and reject unexpected fields or embedded instructions.

4

Require approval tokens for write, send and status-change actions.

5

Use correlation and idempotency keys to link events and prevent duplicate effects.

6

Log tool name, parameters, result, latency, error and approving role without exposing unnecessary data.

In Practice — Integrating Agents with HR Systems and Tools

WORKED EXAMPLE

- ReadPolicy accepts Policy_ID and As_Of_Date and returns approved text, owner and effective date.
- CreateOnboardingTask accepts a validated case key and approval token but cannot edit employee records.
- SendMessage remains disabled in the sandbox; the lab records a simulated message and complete audit line.

DECISION GUIDE

- **Use when**
 - A stable API, connector or controlled file interface has an accountable system owner.
 - The workflow can prove the identity, scope and effect of each tool call.
- **Avoid when**
 - Using a shared administrator account or exposing an entire system to simplify a prototype.
 - Allowing retrieved documents to redefine tool permissions or bypass approval.

Quality Lens — Integrating Agents with HR Systems and Tools

Authorised

- Identity and operation match the current user's and workflow's authority.

Validated

- Schemas, values, destinations and effects are checked.

Auditable

- Every call and approval is linked to one traceable case.

Governance, Fairness and Human Oversight

WHAT

- HR AI governance assigns decision rights and controls across purpose, data, model, tools, users, evaluation, operation and incident response. Human oversight is a designed decision point with evidence, authority, time and fallback, not a generic instruction to review.

WHY IT MATTERS

- HR systems can affect access to work, employee experience and personal data. Trust depends on job-related criteria, proportionate data use, transparent limitations, tested performance, meaningful review and a route for correction or challenge.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Governance, Fairness and Human Oversight

1

Inventory the use case, affected people, decisions, data, vendors and owners.

2

Map benefits and harms, including privacy, bias, exclusion, security and automation overreach.

3

Choose human involvement by impact, uncertainty, reversibility and affected-person expectations.

4

Test source accuracy, unsafe actions, subgroup effects, deferral and override behaviour.

5

Communicate purpose, limitations, data use and routes for questions or correction.

6

Monitor incidents, complaints, drift, reviewer changes and policy updates after launch.

In Practice — Governance, Fairness and Human Oversight

WORKED EXAMPLE

- The candidate assistant may propose rubric evidence but cannot reject, rank on protected traits or contact candidates.
- A trained panel sees the source, unknowns and proposed score and records the final decision and reason.
- Monthly review compares subgroup selection, deferral and correction rates and investigates material gaps.

DECISION GUIDE

- **Use when**
 - Before pilot, after material changes and throughout operation.
 - The organisation can assign business, HR, data, technology, privacy and risk ownership.
- **Avoid when**
 - Treating policy approval as a substitute for testing actual cases and workflow behaviour.
 - Using a human reviewer who lacks evidence, authority, time or a meaningful ability to change the result.

Quality Lens — Governance, Fairness and Human Oversight

Human-centric

- People outcomes, recourse and responsible decision rights shape the design.

Fair

- Job-related criteria and subgroup effects are examined with context.

Transparent

- Purpose, limits, sources, decisions and correction routes are visible.

Rolling Out Agentic AI in Your Organisation

WHAT

- Rollout is a staged operating change. A use case moves through discovery, synthetic sandbox, shadow run, read-only pilot and controlled operation only when value, quality, safety, ownership, support and fallback evidence meet defined gates.

WHY IT MATTERS

- A successful prompt is not an operating service. HR adoption also requires current data, trained reviewers, employee communication, monitoring, support, incident response, change control and a manual route when the system is unavailable or unsafe.

HUMAN CHECK

- Ground the task in approved facts.
- Review evidence, privacy, rights and tone.
- Approve, edit or stop before use.

How It Works — Rolling Out Agentic AI in Your Organisation

1

Prioritise one bounded use case by value, feasibility, data readiness and residual risk.

2

Set a manual baseline, outcome metric, quality metric, risk metric and stop threshold.

3

Test on representative synthetic cases and repair failures before connecting live data.

4

Run in shadow or read-only mode and compare outputs with the approved current process.

5

Pilot with limited users, data and actions; train reviewers and communicate limitations.

6

Monitor, review, version, roll back or retire through named governance and support owners.

In Practice — Rolling Out Agentic AI in Your Organisation

WORKED EXAMPLE

- Asteron pilots policy-grounded onboarding support for one office and two trained HR reviewers.
- The pilot tracks resolution quality, grounding-defect rate, escalation accuracy, correction rate and cycle time.
- Any policy-source conflict or unsafe-action attempt stops the run and triggers the documented manual process.

DECISION GUIDE

- **Use when**
 - The prototype has a measurable baseline, representative tests, named owners and reliable fallback.
 - Users and affected people can understand the system's role and correction path.
- **Avoid when**
 - Scaling because the demonstration looked fluent or because a platform is already licensed.
 - Removing the manual route before reliability, support and recovery have been demonstrated.

Quality Lens — Rolling Out Agentic AI in Your Organisation

Staged

- Exposure grows only after evidence-based gates are met.

Measured

- Value, quality, fairness, safety and operating health are monitored.

Operated

- Owners, support, incidents, changes, fallback and retirement are defined.

Hands-On Labs — Deploy Agentic AI Across the HR Function

Lab 5

- Simulate the Controlled Onboarding Workflow

Lab 6

- Run the Governance Gate and Rollout Plan

Simulate the Controlled Onboarding Workflow

LAB 5

You will convert the Lab 4 onboarding plan into a state-based workflow. ReadPolicy, DraftPlan and CreateTask are represented by explicit tool contracts; write and send effects remain simulated. You will run normal, missing-data, duplicate and policy-conflict cases and reconcile every final state.

You'll build

03-deployment/onboarding-workflow-spec.md and **03-deployment/workflow-run-log.csv** with states, transition rules, tool contracts, four complete traces, human approvals, stop reasons and manual fallback.

1 Map states

2 Contract tools

3 Run four cases

4 Reconcile evidence

Tools: Text editor - spreadsheet - approved AI assistant - workflow-cases.csv - workflow-run-log-template.csv

Simulate the Controlled Onboarding Workflow

Test it

The workflow specification must contain all nine states, three tool contracts, an approval gate and manual fallback. The run log must show WF-001 CLOSED and WF-002 to WF-004 STOPPED with the specified reasons; four terminal states must reconcile, no simulated Task_ID may repeat, and no external action may occur.

Run the Governance Gate and Rollout Plan

LAB 6

You will evaluate the recruitment and onboarding agents against normal, boundary, missing-source, unsafe-action and fairness-relevant cases. You will calculate quality and subgroup indicators, inspect reviewer changes, set release thresholds and write a staged rollout with named owners, monitoring and rollback.

You'll build

03-deployment/governance-evaluation.csv and 03-deployment/hr-agent-rollout-plan.md with risk controls, calculated metrics, fairness investigation, release decision, pilot stages, owners, monitoring and fallback.

1 Map risks

2 Run evaluation

3 Set release gate

4 Plan staged rollout

Tools: Spreadsheet - text editor - governance-test-cases.csv - outputs from Labs 2-5

Run the Governance Gate and Rollout Plan

Test it

The evaluation file must contain ten tests with expected, actual, pass, severity and evidence; all four quality metrics and both subgroup selection and deferral rates must show formulas and denominators. The rollout plan must cover six risk types, six owner roles, four stages, five monitoring dimensions, explicit gates, a release decision and a tested manual fallback. Critical pass rate must be 100% before any pilot approval.

Recap — Deploy Agentic AI Across the HR Function

1

You can now: LO3: build and test a multi-step HR workflow with approved data, limited tools, human gates and complete run evidence

2

You can now: LO4: evaluate privacy, fairness, security and operating risk and produce a staged rollout with owners, metrics and fallback



WRAP-UP

Course Summary & Next Steps

What You Achieved

1

Prompt & Ground

Use an approved AI assistant with a structured prompt, source boundary and human review method for everyday HR work.

2

Design HR Agents

Design bounded recruitment, onboarding, employee-support and communication agents with job-related criteria and explicit human decisions.

3

Automate & Verify

Build and test a multi-step HR workflow that connects approved data and tools while preserving exceptions, approvals and run evidence.

4

Govern & Roll Out

Evaluate privacy, fairness, security and operating risk, then produce a staged rollout plan with owners, metrics, monitoring and fallback.

Continue with a Read-Only HR Pilot

1

Choose one bounded HR support task with current approved sources and an accountable owner.

2

Create a prompt contract, workflow map and representative evaluation set.

3

Run in shadow mode, compare with the approved process and capture every correction.

4

Expand only when value, quality, fairness, privacy, support and fallback thresholds are met.

Keep Practising After the Course

1

Re-run the labs on your own machine — repetition is what makes the skills stick.

2

Apply the techniques to a real process or project in your own organisation.

3

Keep your downloaded C820 course files and use the lab checkpoints to restart safely.

4

Explore the follow-on courses at www.tertiarycourses.com.sg.

C820 - KEEP PEOPLE DECISIONS HUMAN-OWNED

Thank You!

You can now prompt, design, test and roll out HR agents that keep evidence visible and consequential people decisions human-owned.